

# BDDP Final Project Report

Vasyl Korzavatykh, ID: 158669

## Description

For the final project I developed a console application simulating the reservation process in the cinema. I used Apache Cassandra for this project.

## Project structure

The project has 5 main files:

1. **docker-compose.yml** – run with *docker compose up -d* to add and start the nodes cas1 and cas2 (from the image cassandra:latest)
2. **utils.py** – file with all the utility functions shared among files described below
3. **main.py** – file that creates initial keyspace creates / re-creates the tables and populates the tables with sample movies / screenings
4. **stress\_tests.py** – file that runs 3 stress tests listed in the project requirements
5. **user\_console.py** – entrypoint for a sample user that can perform any of the actions described in the section below

## Actions

Sample user can:

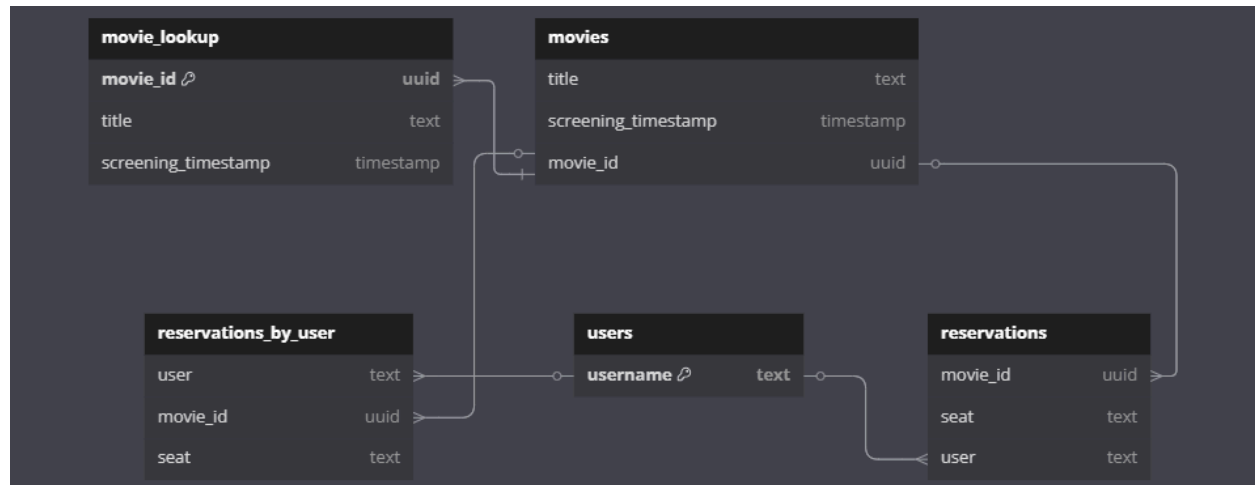
- Make a reservation (choose the movie, screening (time), and a seat)
- View all reservations (made by all the users in the system)
- View their reservations
- Cancel their reservation

## Technical setup

- There are 2 nodes working (cas1 and cas2)
- System functions on any node (the setup is made using `cassandra_driver`)
- The replication factor is set to 2 in the `main.py`
- If one node is down then the system stays functionable, however, operations that require a quorum can't be executed e.g., one can view reservations, however, they can't create one because operations of writing are partly restricted in that scenario

# Database schema

The database schema is (partly) shown below (partly because [dbdiagram.io](https://dbdiagram.io) wouldn't understand Cassandra's primary key notation). Interactive diagram with comments is available via the link: <https://dbdiagram.io/d/cassandra-cinema-reservation-6842fde9a845fcbe14b011b4>



The database has 3 main tables and 2 supplementary ones. The main tables are:

- **users** – table with all the usernames (system is simplistic so there are no passwords)
- **movies** – table where each screening is saved (e.g., movie Inception at 17:00 on June 16, 2025)
- **reservations** – table of reservations where each reservation has corresponding **movie\_id**, seat number, and user name of a person making the reservation

Supplementary tables are added to improve the efficiency of the database and not to use ALLOW FILTERING or indices. There are two of them:

- **reservations\_by\_user** – the user is a partition key while **movie\_id** and seat are clustering keys. This table helps to find all the movies for a given user
- **movie\_lookup** – this table allows to get title and time of a screening knowing the **movie\_id** (primary key)

## Problems encountered

I can't pinpoint any of the serious problems encountered during the project completion, however, what raised a question was a situation when I tried to use 3 nodes with replication factor staying at 2. When I stopped one of the nodes I supposed that the system would stay fully operational but it seems that the node I turned down was the one keeping the replicas and those replicas won't move to the remaining nodes even after some waiting time. I discovered that operations requiring SERIAL consistency wouldn't work because quorum for Paxos wasn't achieved. I believe this is in line with Cassandra functionality, however, not fully intuitive.